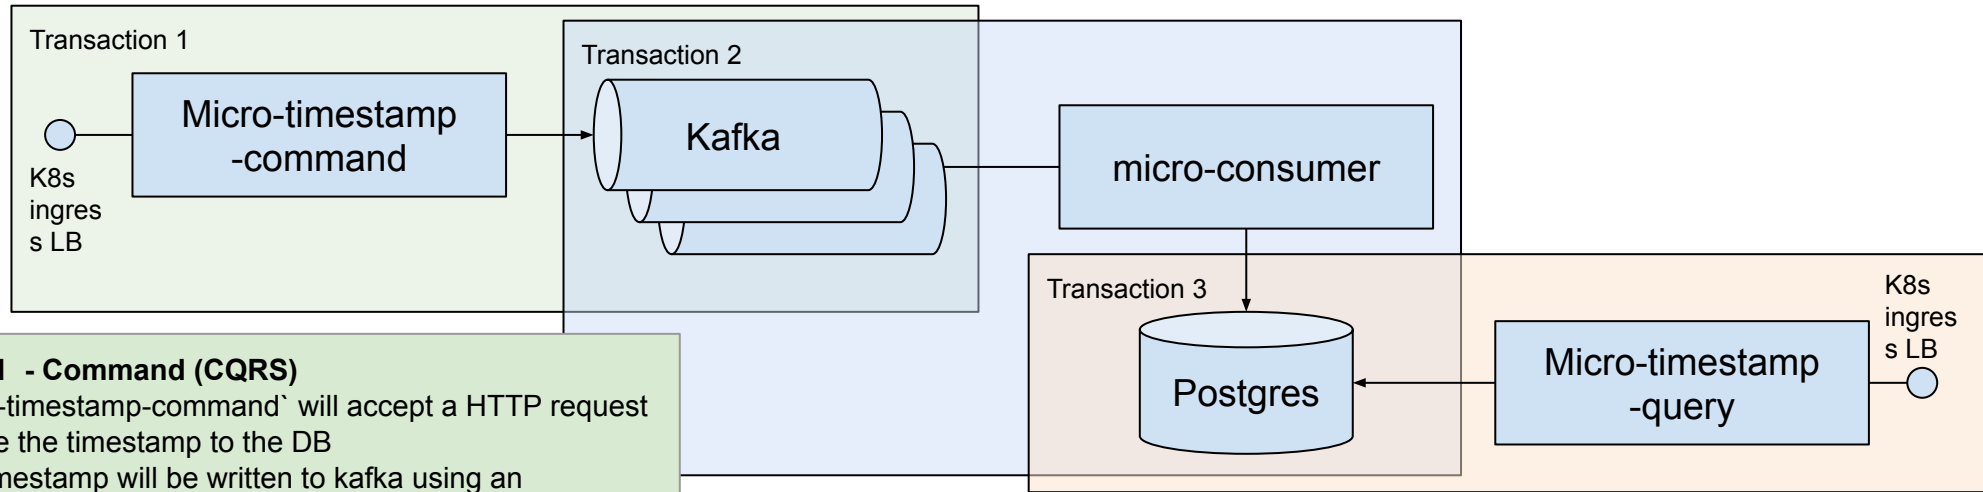


CQRS Implementation



Transaction 1 - Command (CQRS)

- 1) `micro-timestamp-command` will accept a HTTP request to write the timestamp to the DB
- 2) The timestamp will be written to kafka using an idempotent producer (the message key will be a UUID v4)
- 3) The response will return the UUIDv4 for the message key back to the client
- 4) As this is a command-service, it would not retrieve records from postgres, a separate query service would be built for that purpose (time permitting)

Non-functional Requirements

- The kafka log is the golden-record, records in the log are immutable and the postgres view is purely a “temporary” representation of the data
- No retry/Deadlettering will be implemented given time constraints, if there’s a database lock or outage the micro-consumer will not be able to aggregate the records and effectively block downstream processing but will recover upon DB recovery
- The DR process for this will be to re-deploy the consumer and re-point the configuration to another DB, the consumer will rebuild the database dynamically from the golden record in kafka.
- Multiple consumers could build multiple DBs from the single kafka log
- The postgres isn’t HA in this POC as it can be dynamically rebuilt on-the-fly it’s an aggregation of the events in Kafka

Transaction 2 - “View” creation

- 1) `micro-consumer` will use a consumer-group to consume the timestamp log records from Kafka
- 2) Upon consumption of those messages will be written to postgres using a DB transaction
 - a) Open message
 - b) Parse contents
 - c) Open DB transaction
 - i) Write message key (UUID) as the PK on the table
 - ii) Write the timestamp as the value
 - d) Commit transaction
 - i) Upon Success
 - (1) Acknowledge kafka message
 - ii) Upon Failure
 - (1) Do not acknowledge (BLOCK)

Transaction 3 - Query (CQRS)

- 1) Micro-timestamp-query will query the database based on UUID to ensure a record has been created for the requested query

Note this may be omitted given time constraints (as it’s not explicitly required given the design)

Implementation Details

- Where prudent, all services and supporting-services will be deployed as containers (A managed k8s i.e. GKE being the preference)
- In "production" a managed-service will be used to operate kafka e.g. Amazon MSK and provide backup/HA